


Cluster and Cloud Computing – Lecture 5
Cloud Computing & Getting to Grips with the
University of Melbourne Research Cloud
Professor Richard O. Sinnott & *Farzad Khodadadi*
rsinnott@unimelb.edu.au

- Cloud benefits
 - Cloud marketing!?
- The various flavours of cloud computing
 - Introduction to #aaS?
- Break
 - Demonstration of University of Melbourne Research Cloud
- Workshops on Thursday/Friday
 - Workshop: Scripting the Cloud and Ansible
 - Pre-recorded, watch videos and Q&A in workshops


Life before cloud computing



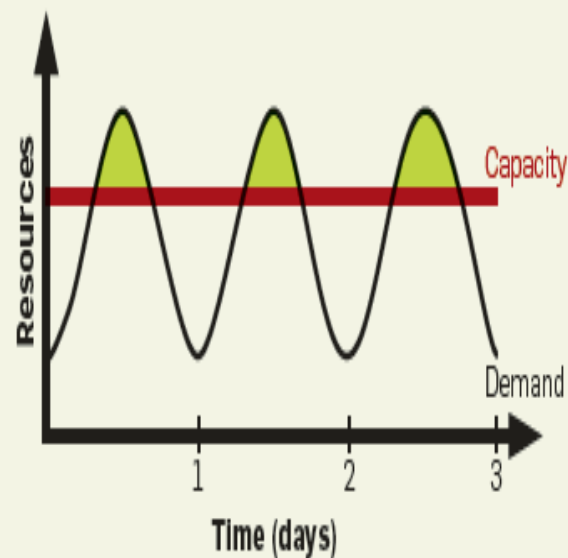
*Capacity vs Utilization curves*⁸

Life before cloud computing...ctd


Figure 2. (a) Even if peak load can be correctly anticipated, without elasticity we waste resources (shaded area) during nonpeak times. (b) Underprovisioning case 1: potential revenue from users not served (shaded area) is sacrificed. (c) Underprovisioning case 2: some users desert the site permanently after experiencing poor service; this attrition and possible negative press result in a permanent loss of a portion of the revenue stream.



(a) Provisioning for peak load



(b) Underprovisioning 1



(c) Underprovisioning 2




Cloud-busting? (~2009)

- To buy or not to buy that is the question...?
 - ScotGrid (www.scotgrid.ac.uk)

- Cost £51

Total CPU time per
between 16 11 2006


18%



Group	Jobs	Kilo SPEC hours	WallClock Hours	CPU Hours	Eff (%)
alice	5	1	1	0	1.0
atlas	466733	2179463	1421698	1050451	74.0
atlasprd	497669	1808452	1179682	1091466	93.0
atlassqm	6420	930	606	465	77.0
babar	2	0	0	0	1.0
biomed	106413	1023678	667761	458138	69.0
biomedsqm	15	0	0	0	5.0
camont	9164	35510	23164	20070	87.0
cms	4857	18557	12105	6686	55.0
cmsprd	677	365	238	138	58.0
dteam	16494	3066	2000	344	17.0
dzero	92	141	92	75	82.0
glbio	33	694	453	446	99.0
glee	58306	2361972	1540752	1527659	99.0
gridpp	55924	1156772	754581	746173	99.0
ilc	18425	144454	94229	68283	72.0
lhcb	52059	956823	624150	584341	94.0
lhcbprd	14468	234657	153070	148710	97.0
lhcbprdm	9653	75513	49258	47224	96.0
ngs	22884	167	109	1	1.0
ops	8577	2573	1678	31	2.0
opssqm	14598	455	296	171	58.0
pheno	178041	1530175	998157	931398	94.0
snemo	14062	94607	61713	13899	24.0
snemosqm	1	0	0	0	0.0
totalcp	15	2	1	0	1.0
zeus	5011	8679	5661	4507	80.0
zeusprdm	64	155	101	0	0.0
Total:	1560662	11637861	7591556	6700676	Overall: 88.0 %

of submitted jobs per selected VO
between 16 11 2006 and 19 3 2009

31%





Cloud-busting? (~2009)

On-Demand Instances

United States

Europe

Standard On-Demand Instances

Linux/UNIX Usage

Windows Usage

1.7Gb memory/160GB disk
7.5Gb memory/850GB disk
15Gb memory/1690GB disk

\$0.11 per hour

\$0.135 per hour

\$0.44 per hour

\$0.54 per hour

\$0.88 per hour

\$1.08 per hour

High CPU On-Demand Instances

Linux/UNIX Usage

Windows Usage

Medium

\$0.22 per hour

\$0.32 per hour

Extra Large

\$0.88 per hour

\$1.28 per hour

Amazon Elastic Block Store

United States

Europe

Amazon EBS Volumes

- \$0.11 per GB-month of provisioned storage
- \$0.11 per 1 million I/O requests

Amazon EBS Snapshots to Amazon S3 (priced the same as Amazon S3)

- \$0.18 per GB-month of data stored
- \$0.012 per 1,000 PUT requests (when saving a snapshot)
- \$0.012 per 10,000 GET requests (when loading a snapshot)

Data Transfer

Internet Data Transfer

The pricing below is based on data transferred "in" and "out" of Amazon EC2.

Data Transfer In

All Data Transfer \$0.10 per GB

Data Transfer Out

First 10 TB per Month \$0.17 per GB

Next 40 TB per Month \$0.13 per GB

Next 100TB per Month \$0.11 per GB

Over 150 TB per Month \$0.10 per GB

- \$1 = £0.69 (back then!)
 - £0.30 * 6,700,676 CPU hours
 - = £2,010,202 for just compute on-demand + data + networking + ...?
- Now...???
 - <https://aws.amazon.com/ec2/pricing/>



AWS EC2 Pricing (just EC2!)




- **On demand**
 - pay for compute capacity used; no long-term commitments; increase/decrease as required
- **Spot pricing**
 - when have flexibility in using resources, e.g. get resources to run my jobs when price drops below...
- **Saving plans**
 - Longer term arrangements, need to use regularly
- **Dedicated Hosts**
 - Physical dedicated servers; can use own licensed software
- **Reserved instances**
 - for applications with known usage patterns, e.g. fixed reserve resources in given availability zone

- NIST definition: “Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.”
 - » National Institute of Standards and Technology
(<http://dx.doi.org/10.6028/NIST.SP.800-145>)
- Focus of today is to get you up and running on the Cloud and explore the technologies related to the underlined
 - Later lecture will do compare/contrast with AWS
 - ...and then how do Clouds actually work (hypervisors etc)



The Most Common Cloud Models

Deployment
Models



Delivery
Models

Software as a
Service (SaaS)

Platform as a
Service (PaaS)

Infrastructure as a
Service (IaaS)

Essential
Characteristics

- On-demand self-service
- Broad network access
- Resource pooling
- Rapid elasticity
- Measured service

Note also XYZ-as-a-Service

- Pros
 - Utility computing
 - Can focus on core business
 - Cost-effective
 - “Right-sizing”
 - Democratisation of computing
- Cons
 - Security
 - Loss of control
 - Possible lock-in
 - Dependency of Cloud provider continued existence



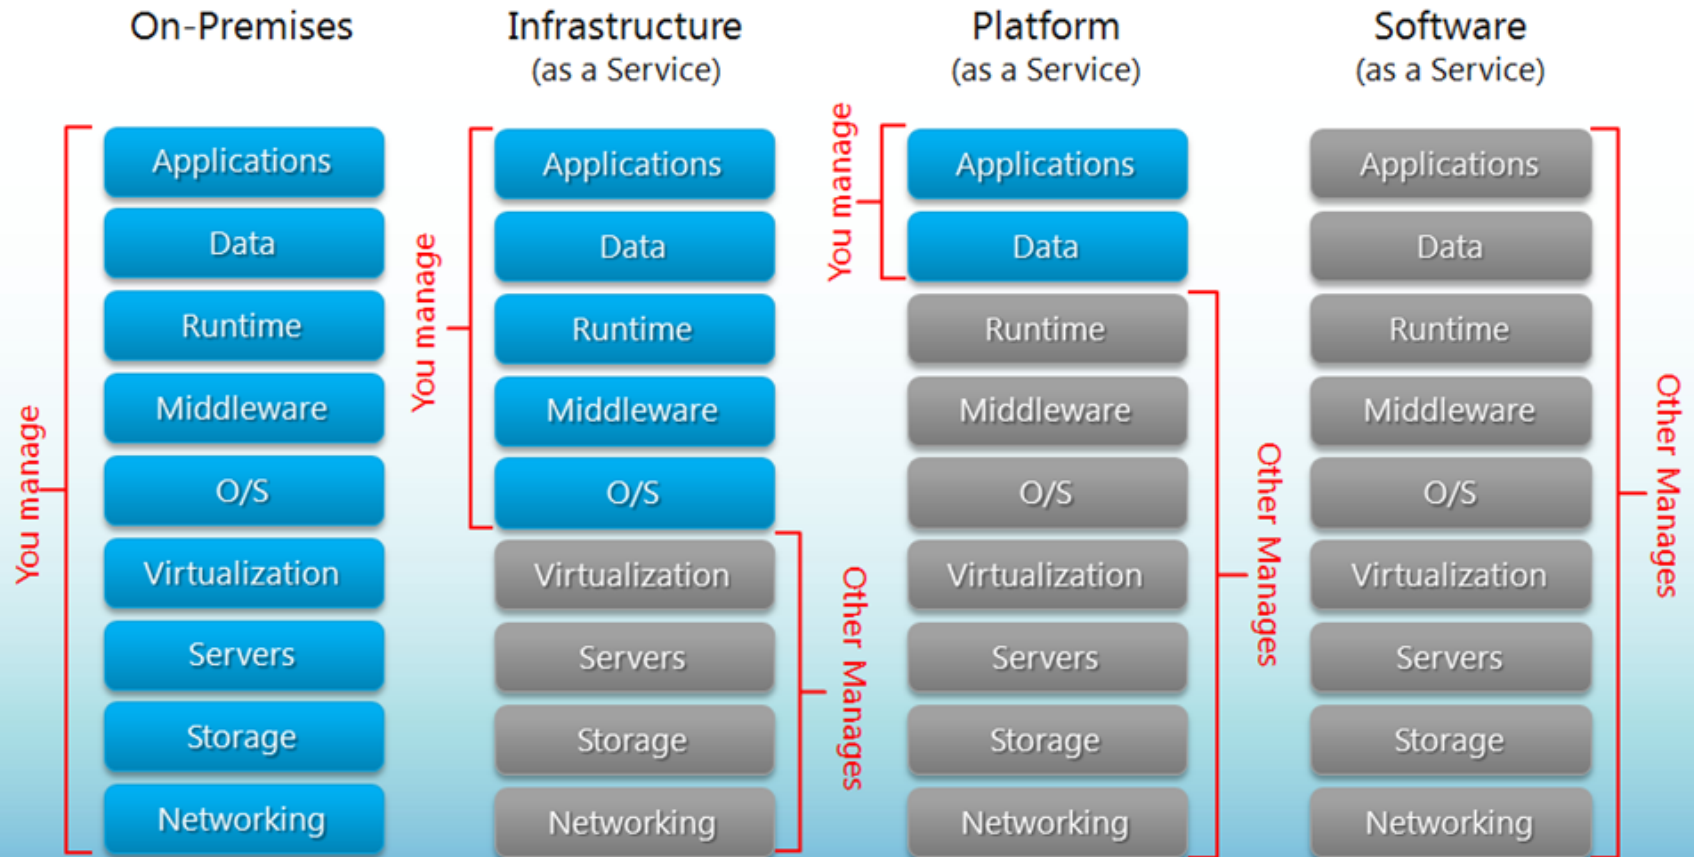
Private Clouds

- Pros
 - Control
 - Consolidation of resources
 - Easier to secure
 - More trust
- Cons
 - Relevance to core business?
 - e.g. Netflix to Amazon
 - Staff/management overheads
 - Hardware obsolescence
 - Over/under utilisation challenges



- Examples
 - Eucalyptus, VMWare vCloud Hybrid Service
- Pros
 - Cloud-bursting
 - Use private cloud, but burst into public cloud when needed
- Cons
 - How do you move data/resources when needed?
 - How to decide (in real time?) what data can go to public cloud?
 - Is the public cloud compliant with PCI-DSS (Payment Card Industry – Data Security Standard)?

Separation of Responsibilities





Public SaaS examples

- Gmail
- Sharepoint
- Salesforce.com CRM
- On-live
- Gaikai
- Microsoft Office 365
- Some definitions include those that do not require payment, e.g. ad-supported sites



Public PaaS Examples


Cloud Name	Language and Developer Tools	Programming Models Supported by Provider	Target Applications and Storage Options
Google App Engine	Python, Java, Go, PHP + JVM languages (scala, groovy, jruby)	MapReduce, Web, DataStore, Storage and other APIs	Web applications and BigTable storage
Salesforce.com's Force.com	Apex, Eclipsed-based IDE, web-based wizard	Workflow, excel-like formula, web programming	Business applications such as CRM
Microsoft Azure	.NET, Visual Studio, Azure tools	Unrestricted model	Enterprise and web apps
Amazon Elastic MapReduce	Hive, Pig, Java, Ruby etc.	MapReduce	Data processing and e-commerce
Aneka	.NET, stand-alone SDK	Threads, task, MapReduce	.NET enterprise applications, HPC



Infrastructure As A Service (IaaS)

- Primary focus of this course...
- Many providers
 - Amazon Web Services (Market leader)
 - <http://aws.amazon.com>
 - Oracle Public Cloud
 - <https://cloud.oracle.com/>
 - Rackspace Cloud
 - www.rackspace.com
 - *CenturyLink, CloudSigma, DigitalOcean, DimensionData, GoGrid, Helio, Internap, Joyent, ProfitBricks, Verizon, ...*
 - NeCTAR/Openstack Research Cloud
 - www.nectar.org.au

Cost Comparison



See also <https://www.infoworld.com/article/3237566/cloud-pricing-comparison-aws-vs-azure-vs-google-vs-ibm.html>



NeCTAR/ARDC/UniMelb Research Cloud

- National eResearch Collaboration Tools and Resources (NeCTAR – www.nectar.org.au)
 - \$50m+\$10m+\$10m+\$72m (Australian Research Data Commons)... federal funding
 - Originally lead by University of Melbourne
 - Had four key strands
 - ~~National Servers Program~~
 - Research Cloud Program
 - OpenStack IaaS
 - 4Gb-64Gb (mostly Linux flavours)
 - 30,000 physical servers available across different availability zones
 - » Being upgraded continually!
 - ~~eResearch Tools Program~~
 - Virtual Laboratories Program
 - Astro,
 - Genomics,
 - Humanities,
 - Climate,
 - Nano-,
 - ...endocrine genomics



- Research Data Services (RDS)
- \$50m+\$10m+\$10m+\$72m project to establish data storage resources across Australia
 - ~100 Petabytes national data storage
 - Victoria Node (VicNode)
 - UniMelb, UniMonash for Vic-wide “nationally significant data sets”
 - Used by many diverse communities



- Petascale Campus Initiative
 - <https://research.unimelb.edu.au/infrastructure/petascale-campus-initiative>
- Research Computing Services
 - <https://research.unimelb.edu.au/infrastructure/research-computing-services>



- Based on OpenStack
 - Open source cloud technology (more later lecture)
- Many associated/underpinning services
 - Compute Service (code-named **Nova**)
 - Image Service (code-named **Glance**)
 - Block Storage Service (code named **Cinder**)
 - Object Storage Service (code-named **Swift**)
 - Security Management (code-named **Keystone**)
 - Orchestration Service (code-named **Heat**)
 - Network Service (code-named **Neutron**)
 - Metering Service (code-named **Ceilometer**)
 - ...

<https://docs.openstack.org/rocky/>



UniMelb Research Cloud



Melbourne Research Cloud / Project / unimelb-comp90024-2020-grp-10

PROJECTS | USER | SUPPORT

Project / Compute / Overview

PROJECT

API Access

Compute

Overview

Instances

Images

Key Pairs

Server Groups

Volumes

Network

Orchestration

DNS

Object Store

Users

APP CATALOG

ALLOCATIONS

Limit Summary

Compute



Instances

Used 0 of 4



VCPUs

Used 0 of 8



RAM

Used 0Bytes of 36GB

Volume



Volumes

Used 0 of 250



Volume Snapshots

Used 0 of 250



Volume Storage

Used 0Bytes of 250GB

Network



Floating IPs

Allocated 0 of 0



Security Groups

Used 3 of 30



Security Group Rules

Used 11 of 150



Networks

Used 0 of 0



0

Ports

Used 0 (No Limit)



Routers

Used 0 of 0



BREAK



UniMelb/NeCTAR Demo (Crib Sheets LMS)

- NeCTAR (<https://dashboard.rc.nectar.org.au>)
- UniMelb (<https://dashboard.cloud.unimelb.edu.au>)
- Note that there are some differences
 - e.g. limited number of IP addresses



UniMelb/NeCTAR Demo (Crib Sheets LMS)

- Dashboard
 - Overview (Resource allocation)
 - Instances (Create, terminate and configure instances)
 - Volumes (Create, terminate, attach and backup)
 - Images (Create image, list image)
 - Access & Security (Security groups, key pairs, API access)
 - Object store (Store data as an object)
- Launching a new VM
 - Flavor (defines the compute, memory, and storage capacity)
 - Ephemeral disk (not persistence storage!!!)
 - Create key pair
 - `ssh-keygen -t rsa -f cloud.key` (Unix/Linux/MacOS, see Putty for Windows)
 - `chmod 600 cloud.key` (or get UNPROTECTED PRIVATE KEY FILE! error)



NeCTAR Demo (Crib Sheets LMS)

- Launching a new VM (continue)
 - Copy pub key (cloud.pub)
 - Select key pair
 - Select security group
 - Availability zone
- Connecting to VM via SSH
 - Private key (your secret key to keep secure)
 - `ssh -i <private-key> <username>@<hostname>`
 - E.g. `ssh -i cloud.key ubuntu@example.com`
- Create a Volume
 - Must be in the same availability zone as the instance

Note: you can't create volumes until
you have some (=assignment 2)!



NeCTAR Demo (Crib Sheets LMS)

- Attach a volume
 - Check the device name: *sudo fdisk -l*
 - *sudo* = runs commands with security privileges of another user
 - (by default: superuser)
 - Create the mounting point: *sudo mkdir /mnt/demo*
 - Format the volume: *sudo mkfs.ext4 /dev/vdb*
 - Mount the volume: *sudo /dev/vdb /mnt/demo*
 - Check the result: *df -h*
- Installing an application
 - *sudo apt-get install vim*
- Create snapshots
 - Snapshot for an instance
 - Snapshot for a volume

Note: if you don't know Linux or what commands like *df -h* do, then google is your friend! See later too...



NeCTAR Demo (Crib Sheets LMS)

- Restore a snapshot
 - Create an instance from an instance snapshot
 - Create a volume from a volume snapshot
- Setting up security groups
 - Security groups act as a virtual firewall that controls the traffic for one or more instances
 - It contains a set of security rules
 - Default security group only allows SSH (from anywhere)
 - Create a security group
 - Create a rule
 - CIDR (https://en.wikipedia.org/wiki/Classless_Inter-Domain_Routing)
 - 128.250.0.0/16 (All IPs from UoM)
 - x.x.x.x/32 (x.x.x.x only, use CIDR calculate for IP range calculation)
 - Security Groups

Note: be careful with security.
Only open ports that are needed!



Basic recipes to follow...

- SSH

- `ssh-keygen -t rsa -f cloud.key` (Unix/Linux/macOS, see Putty for Windows)
- `chmod 600 cloud.key` (or get **UNPROTECTED PRIVATE KEY FILE!** error)
- `ssh -i cloud.key ubuntu@<instance_ip>` (Ubuntu)
- `ssh -i cloud.key ec2_user@<instance_ip>` (Amazon Linux, RHEL)

- Volume

- `sudo fdisk -l`
 - `sudo` = runs commands with security privileges of another user (by default: superuser)
- `sudo mkdir /mnt/demo`
- `sudo mkfs.ext4 /dev/vdb`
 - `mkfs` = make file system
 - `ext4` = type of file system (ext2, ext3, ext4)
 - See <http://www.thegeekstuff.com/2011/05/ext2-ext3-ext4/> for details on differences
- `sudo mount /dev/vdb /mnt/demo`
- `df -h`



Basic recipes to follow...

- Install software
 - *sudo apt-get install vim*
 - *apt-get* installs/removes packages on Ubuntu installations
- Common Shell Commands
 - <http://www.dummies.com/computers/operating-systems/linux/common-linux-commands/>