Workshop 9 - Big Data Analytics

Luca Morandini
Data Architect – AURIN Project
University of Melbourne
luca.morandini@unimelb.edu.au

Starting Spark With Docker

- The instructions on how to build a mini-cluster (a Spark master and one Spark worker Docker containers), and an example of how to run a word-count MapReduce job is at the usual **https://github.com/AURIN/comp90024** repository, go under the **spark** directory and follow the README.

Visualizing a Spark Cluster

- Spark *Web-UI* allows to graphically see nodes (port 8080)

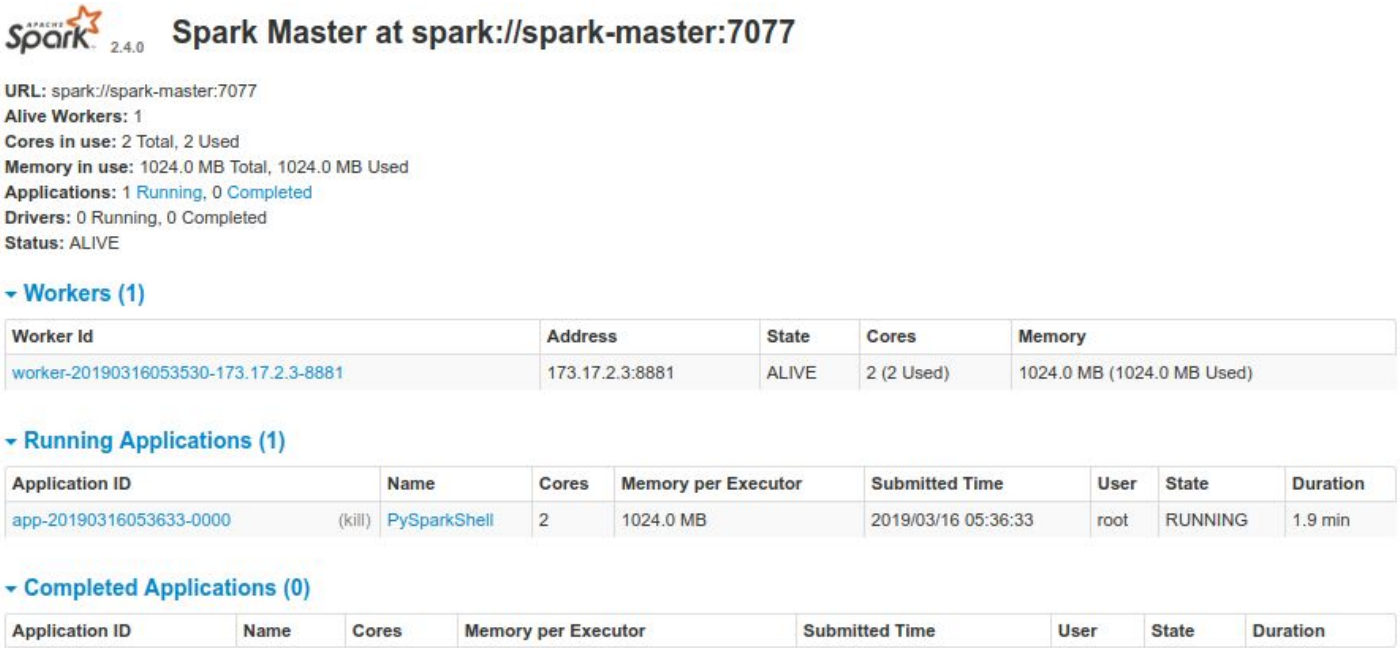

Spark 2.4.0 **Spark Master at spark://spark-master:7077**

URL: spark://spark-master:7077
Alive Workers: 1
Cores in use: 2 Total, 2 Used
Memory in use: 1024.0 MB Total, 1024.0 MB Used
Applications: 1 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190316053530-173.17.2.3-8881	173.17.2.3:8881	ALIVE	2 (2 Used)	1024.0 MB (1024.0 MB Used)

Running Applications (1)

Application ID		Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
app-20190316053633-0000	(kill)	PySparkShell	2	1024.0 MB	2019/03/16 05:36:33	root	RUNNING	1.9 min

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Visualizing Spark Jobs: Executors

- Spark *Web-UI* allows as well to see the division of jobs in stages and tasks, and the allocation of executors to nodes (port 4040)

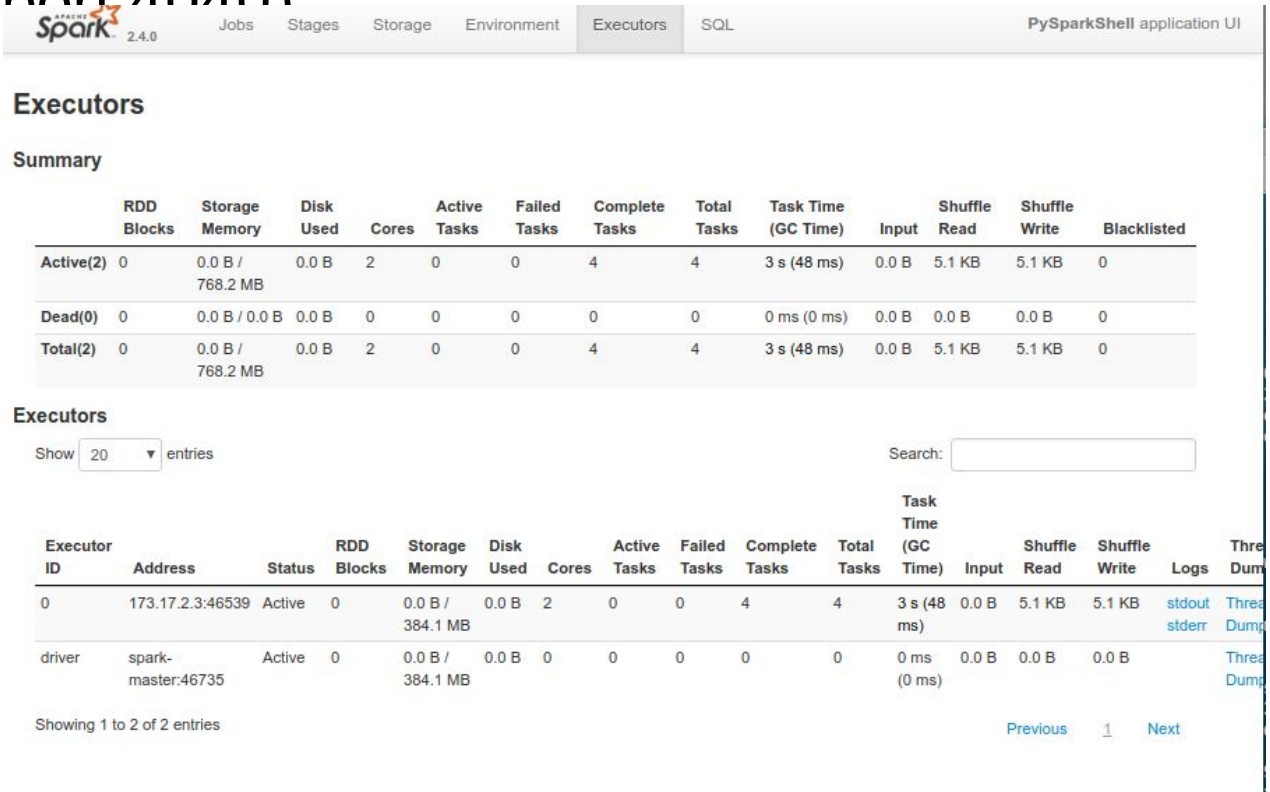

Visualizing Spark Jobs: Jobs

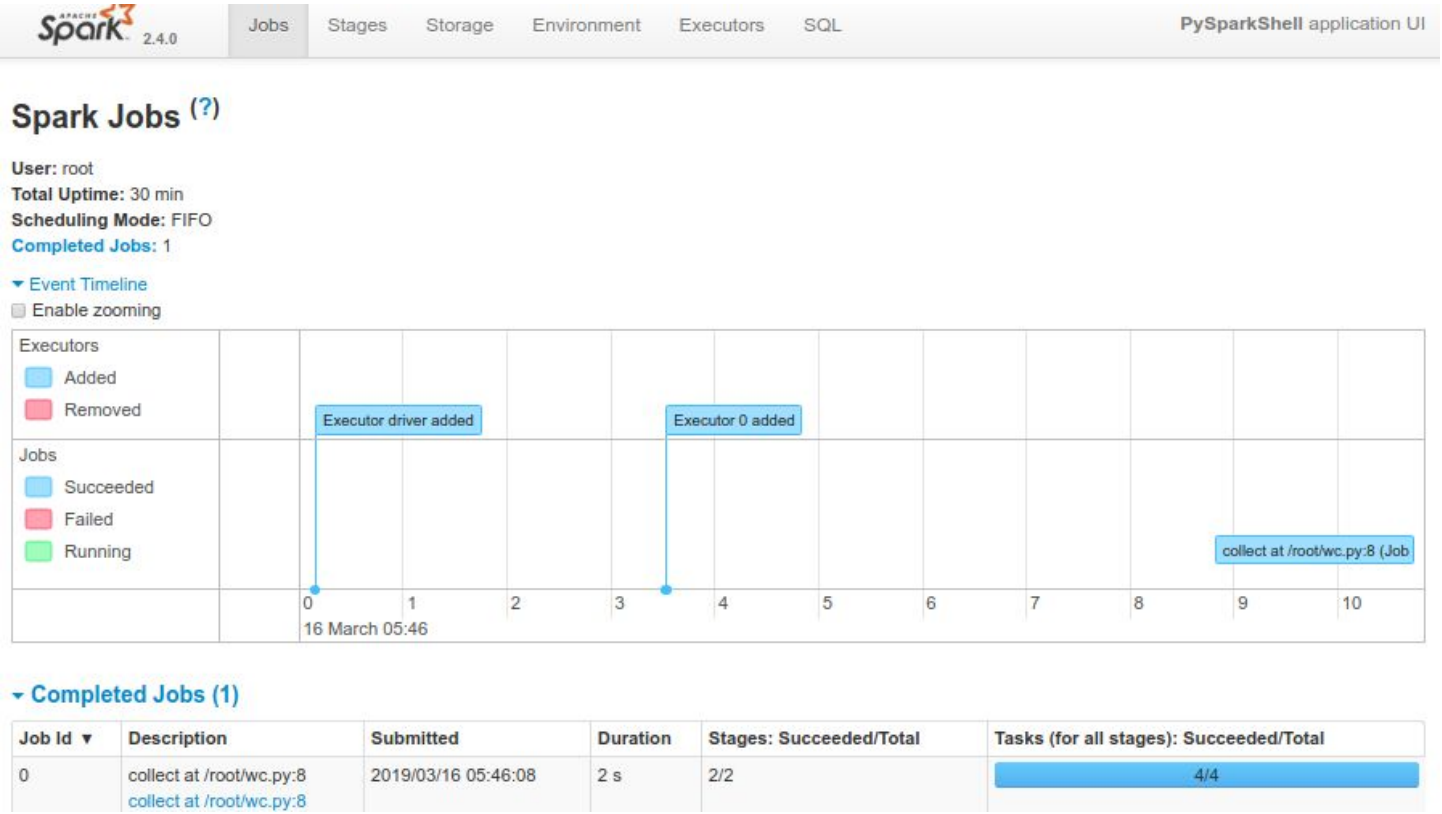

Visualizing Spark Jobs: Single Job

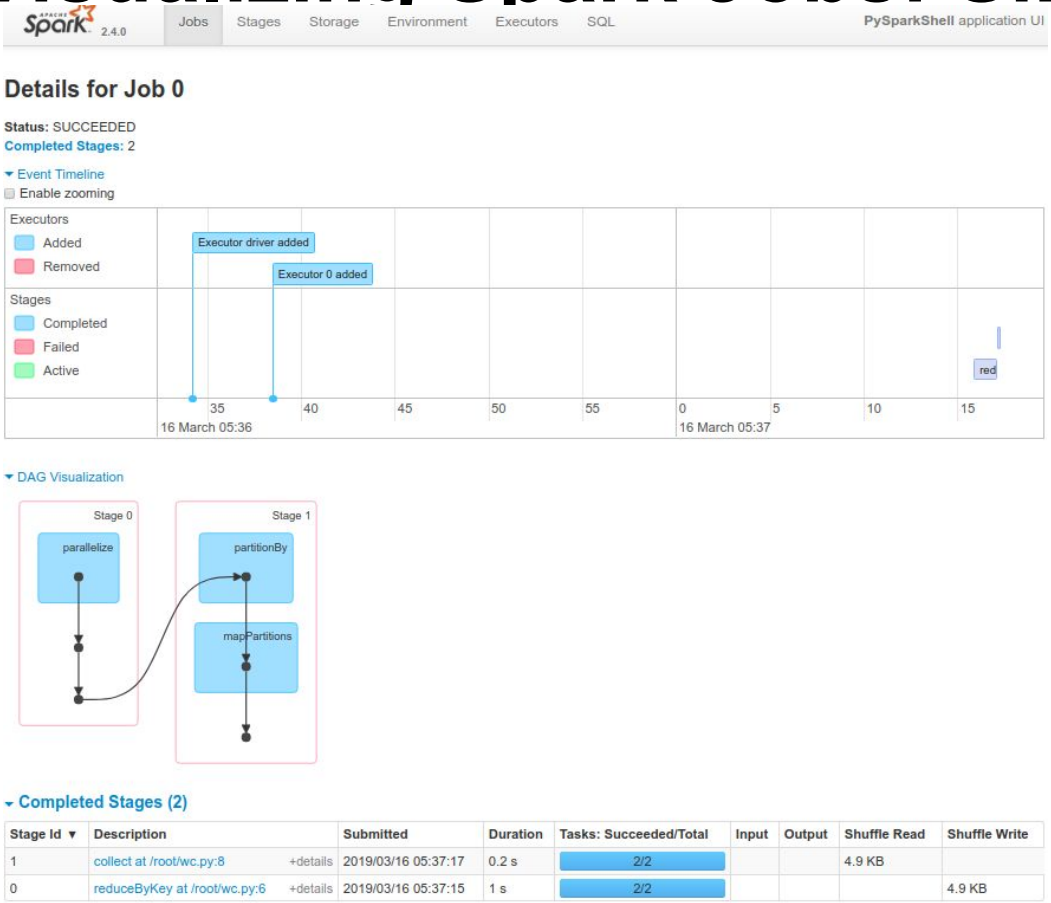

Visualizing Spark Jobs: Tasks of a Stage